

# COMS 4721: Machine Learning for Data Science

## Lecture 18, 4/9/2019

Prof. John Paisley

Department of Electrical Engineering  
& Data Science Institute  
Columbia University

# COLLABORATIVE FILTERING

# OBJECT RECOMMENDATION

Matching consumers to products is an important practical problem.

We can often make these connections using user feedback about subsets of products. To give some prominent examples:

- ▶ Netflix lets users rate movies
- ▶ Amazon lets users rate products and write reviews about them
- ▶ Yelp lets users rate businesses, write reviews, upload pictures
- ▶ YouTube lets users like/dislike videos and write comments

Recommendation systems use this information to help recommend new things to customers that they may like.

# CONTENT FILTERING

One strategy for object recommendation is:

**Content filtering:** Use known information about the products and users to make recommendations. Create profiles based on

- ▶ Products: movie information, price information, product descriptions
- ▶ Users: demographic information, questionnaire information

**Example:** A fairly well known example is the online radio Pandora, which uses the “Music Genome Project.”

- ▶ An expert scores a song based on hundreds of characteristics
- ▶ A user also provides information about his/her music preferences
- ▶ Recommendations are made based on pairing these two sources

# COLLABORATIVE FILTERING

Content filtering requires a lot of information that can be difficult and expensive to collect. Another strategy for object recommendation is:

**Collaborative filtering (CF):** Use previous users' input/behavior to make future recommendations. Ignore any *a priori* user or object information.

- ▶ CF uses the ratings of similar users to predict my rating.
- ▶ CF is a domain-free approach. It doesn't need to know what is being rated, just who rated what, and what the rating was.

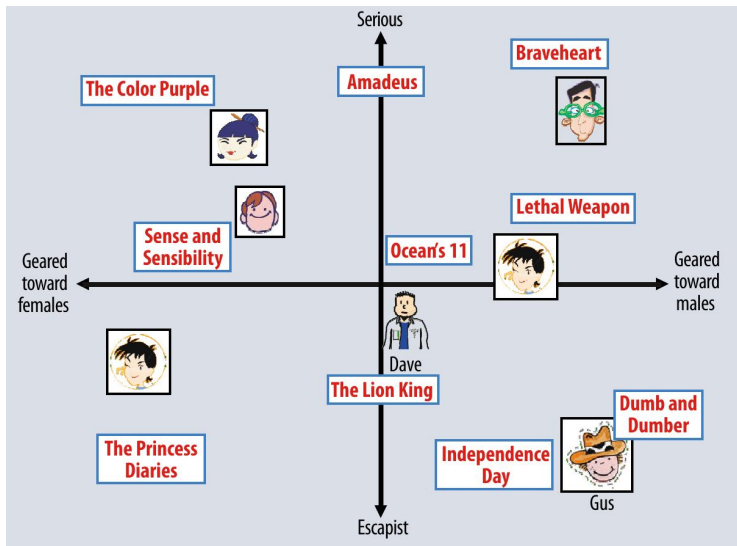
One CF method uses a *neighborhood-based* approach. For example,

1. define a similarity score between me and other users based on how much our overlapping ratings agree, then
2. based on these scores, let others "vote" on what I would like.

These filtering approaches are not mutually exclusive. Content information can be built into a collaborative filtering system to improve performance.

# LOCATION-BASED CF METHODS (INTUITION)

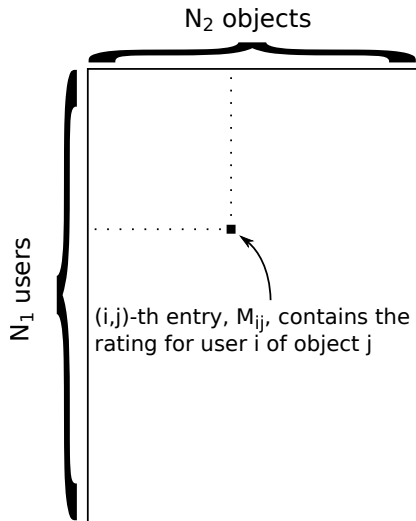
*Location-based* approaches embed users and objects into points in  $\mathbb{R}^d$ .



<sup>1</sup> Koren, Y., Robert B., and Volinsky, C.. "Matrix factorization techniques for recommender systems." Computer 42.8 (2009): 30-37.

# MATRIX FACTORIZATION

# MATRIX FACTORIZATION



Matrix factorization (MF) gives a way to learn user and object locations.

First, form the rating matrix  $M$ :

- ▶ Contains every user/object pair.
- ▶ Will have many missing values.
- ▶ The goal is to fill in these missing values.

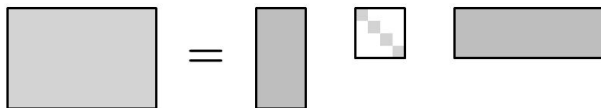
MF and recommendation systems:

- ▶ We have prediction of every missing rating for user  $i$ .
- ▶ Recommend the highly rated objects among the predictions.



# SINGULAR VALUE DECOMPOSITION

Our goal is to factorize the matrix  $M$ . We've discussed one method already.

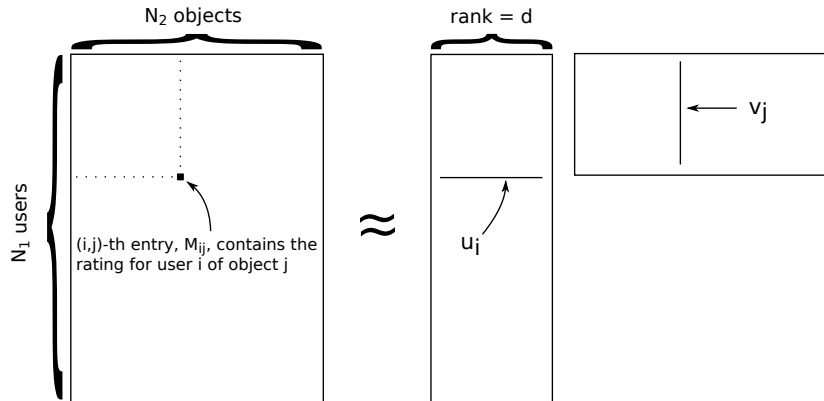

$$\begin{matrix} \mathbf{M} & = & \mathbf{U} & \mathbf{S} & \mathbf{V}^T \\ (n \times d) & & (n \times r) & (r \times r) & (r \times d) \end{matrix}$$

**Singular value decomposition:** Every matrix  $M$  can be written as  $M = USV^T$ , where  $U^T U = I$ ,  $V^T V = I$  and  $S$  is diagonal with  $S_{ii} \geq 0$ .

$r = \text{rank}(M)$ . When it's small,  $M$  has fewer “degrees of freedom.”

Collaborative filtering with matrix factorization is intuitively similar.

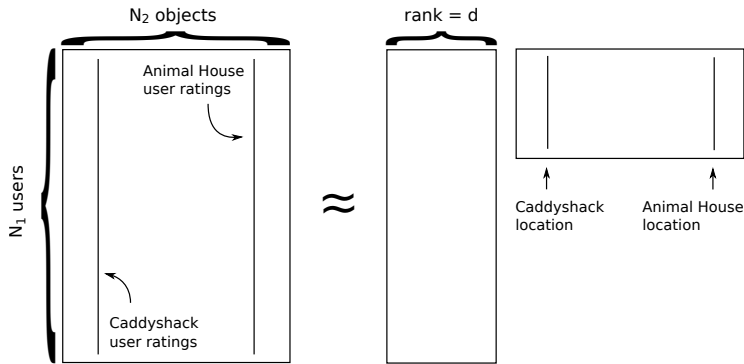
# MATRIX FACTORIZATION



We will define a model for learning a low-rank factorization of  $M$ . It should:

1. Account for the fact that most values in  $M$  are missing
2. Be low-rank, where  $d \ll \min\{N_1, N_2\}$  (e.g.,  $d \approx 10$ )
3. Learn a location  $u_i \in \mathbb{R}^d$  for user  $i$  and  $v_j \in \mathbb{R}^d$  for object  $j$

# LOW-RANK MATRIX FACTORIZATION

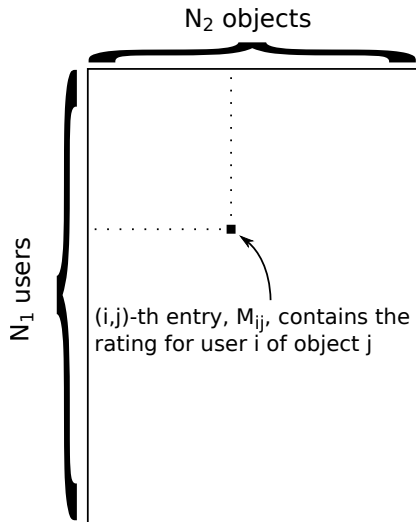


Why learn a low-rank matrix?

- ▶ We think that many columns should look similar. For example, movies like *Caddyshack* and *Animal House* should have **correlated** ratings.
- ▶ Low-rank means that the  $N_1$ -dimensional columns don't "fill up"  $\mathbb{R}^{N_1}$ .
- ▶ Since  $> 95\%$  of values may be missing, a low-rank restriction gives hope for filling in missing data because it models correlations.

# PROBABILISTIC MATRIX FACTORIZATION

# SOME NOTATION



- Let the set  $\Omega$  contain the pairs  $(i,j)$  that are observed. In other words,

$$\Omega = \{(i,j) : M_{ij} \text{ is measured}\}.$$

So  $(i,j) \in \Omega$  if user  $i$  rated object  $j$ .

- Let  $\Omega_{u_i}$  be the index set of objects rated by user  $i$ .
- Let  $\Omega_{v_j}$  be the index set of users who rated object  $j$ .

# PROBABILISTIC MATRIX FACTORIZATION

## Generative model

For  $N_1$  users and  $N_2$  objects, generate

**User locations:**  $u_i \sim N(0, \lambda^{-1}I), \quad i = 1, \dots, N_1$

**Object locations:**  $v_j \sim N(0, \lambda^{-1}I), \quad j = 1, \dots, N_2$

Given these locations the distribution on the data is

$$M_{ij} \sim N(u_i^T v_j, \sigma^2), \quad \text{for each } (i, j) \in \Omega.$$

Comments:

- ▶ Since  $M_{ij}$  is a rating, the Gaussian assumption is clearly wrong.
- ▶ However, the Gaussian is a convenient assumption. The algorithm will be easy to implement, and the model works well.

# MODEL INFERENCE

**Q:** There are many missing values in the matrix  $M$ . Do we need some sort of EM algorithm to learn all the  $u$ 's and  $v$ 's?

- ▶ Let  $M_o$  be the part of  $M$  that is observed and  $M_m$  the missing part. Then

$$p(M_o|U, V) = \int p(M_o, M_m|U, V) dM_m.$$

- ▶ Recall that EM is a **tool** for maximizing  $p(M_o|U, V)$  over  $U$  and  $V$ .
- ▶ Therefore, it is only needed when
  1.  $p(M_o|U, V)$  is hard to maximize,
  2.  $p(M_o, M_m|U, V)$  is easy to work with, and
  3. the posterior  $p(M_m|M_o, U, V)$  is known.

**A:** If  $p(M_o|U, V)$  doesn't present any problems for inference, then no.

(Similar conclusion in our MAP scenario, maximizing  $p(M_o, U, V)$ .)

# MODEL INFERENCE

To test how hard it is to maximize  $p(M_o, U, V)$  over  $U$  and  $V$ , we have to

1. Write out the joint likelihood
2. Take its natural logarithm
3. Take derivatives with respect to  $u_i$  and  $v_j$  and see if we can solve

The joint likelihood of  $p(M_o, U, V)$  can be factorized as follows:

$$p(M_o, U, V) = \underbrace{\left[ \prod_{(i,j) \in \Omega} p(M_{ij} | u_i, v_j) \right]}_{\text{conditionally independent likelihood}} \times \underbrace{\left[ \prod_{i=1}^{N_1} p(u_i) \right] \left[ \prod_{j=1}^{N_2} p(v_j) \right]}_{\text{independent priors}}.$$

By definition of the model, we can write out each of these distributions.



# MAXIMUM A POSTERIORI

## Log joint likelihood and MAP

The MAP solution for  $U$  and  $V$  is the maximum of the log joint likelihood

$$U_{\text{MAP}}, V_{\text{MAP}} = \arg \max_{U, V} \sum_{(i,j) \in \Omega} \ln p(M_{ij} | u_i, v_j) + \sum_{i=1}^{N_1} \ln p(u_i) + \sum_{j=1}^{N_2} \ln p(v_j)$$

Calling the MAP objective function  $\mathcal{L}$ , we want to maximize

$$\mathcal{L} = - \sum_{(i,j) \in \Omega} \frac{1}{2\sigma^2} \|M_{ij} - u_i^T v_j\|^2 - \sum_{i=1}^{N_1} \frac{\lambda}{2} \|u_i\|^2 - \sum_{j=1}^{N_2} \frac{\lambda}{2} \|v_j\|^2 + \text{constant}$$

The squared terms appear because all distributions are Gaussian.

# MAXIMUM A POSTERIORI

To update each  $u_i$  and  $v_j$ , we take the derivative of  $\mathcal{L}$  and set to zero.

$$\nabla_{u_i} \mathcal{L} = \sum_{j \in \Omega_{u_i}} \frac{1}{\sigma^2} (M_{ij} - u_i^T v_j) v_j - \lambda u_i = 0$$

$$\nabla_{v_j} \mathcal{L} = \sum_{i \in \Omega_{v_j}} \frac{1}{\sigma^2} (M_{ij} - v_j^T u_i) u_i - \lambda v_j = 0$$

We can solve for each  $u_i$  and  $v_j$  individually (therefore EM isn't required),

$$u_i = \left( \lambda \sigma^2 I + \sum_{j \in \Omega_{u_i}} v_j v_j^T \right)^{-1} \left( \sum_{j \in \Omega_{u_i}} M_{ij} v_j \right)$$

$$v_j = \left( \lambda \sigma^2 I + \sum_{i \in \Omega_{v_j}} u_i u_i^T \right)^{-1} \left( \sum_{i \in \Omega_{v_j}} M_{ij} u_i \right)$$

However, we can't solve for all  $u_i$  and  $v_j$  at once to find the MAP solution. Thus, as with K-means and the GMM, we use a coordinate ascent algorithm.

# PROBABILISTIC MATRIX FACTORIZATION

## MAP inference coordinate ascent algorithm

---

**Input:** An incomplete ratings matrix  $M$ , as indexed by the set  $\Omega$ . Rank  $d$ .

**Output:**  $N_1$  user locations,  $u_i \in \mathbb{R}^d$ , and  $N_2$  object locations,  $v_j \in \mathbb{R}^d$ .

**Initialize** each  $v_j$ . For example, generate  $v_j \sim N(0, \lambda^{-1}I)$ .

**for** each iteration **do**

**for**  $i = 1, \dots, N_1$  **update user location**

$$u_i = \left( \lambda \sigma^2 I + \sum_{j \in \Omega_{u_i}} v_j v_j^T \right)^{-1} \left( \sum_{j \in \Omega_{u_i}} M_{ij} v_j \right)$$

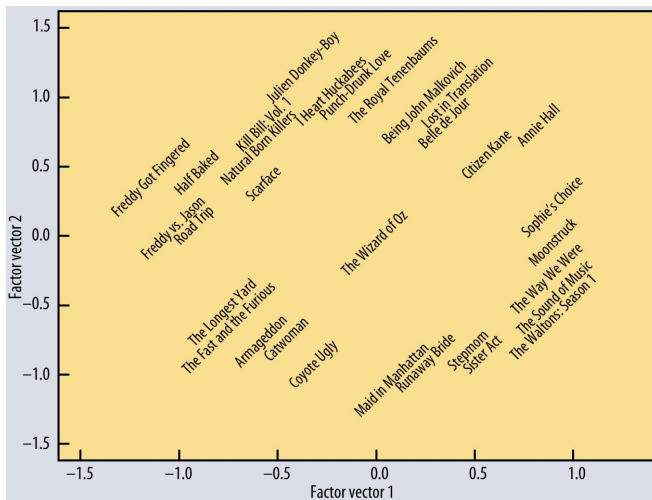
**for**  $j = 1, \dots, N_2$  **update object location**

$$v_j = \left( \lambda \sigma^2 I + \sum_{i \in \Omega_{v_j}} u_i u_i^T \right)^{-1} \left( \sum_{i \in \Omega_{v_j}} M_{ij} u_i \right)$$

**Predict** that user  $i$  rates object  $j$  as  $u_i^T v_j$  rounded to closest rating option

---

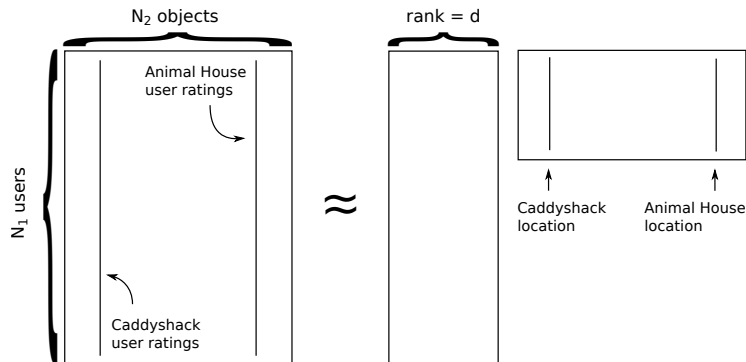
# ALGORITHM OUTPUT FOR MOVIES



Hard to show in  $\mathbb{R}^2$ , but we get locations for movies and users. Their relative locations captures relationships (that can be hard to explicitly decipher).

<sup>1</sup> Koren, Y., Robert B., and Volinsky, C.. "Matrix factorization techniques for recommender systems." Computer 42.8 (2009): 30-37.

# ALGORITHM OUTPUT FOR MOVIES

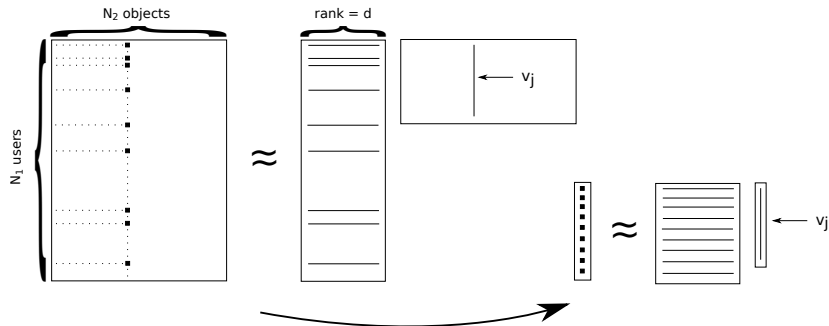


Returning to *Animal House* ( $j$ ) and *Caddyshack* ( $j'$ ), it's easy to understand the relationship between their locations  $v_j$  and  $v_{j'}$ :

- ▶ For these two movies to have similar rating patterns, their respective  $v$ 's must be similar (i.e., close to each other in  $\mathbb{R}^d$ ).
- ▶ The same holds for users who have similar tastes across movies.

# MATRIX FACTORIZATION AND RIDGE REGRESSION

# MATRIX FACTORIZATION AND RIDGE REGRESSION



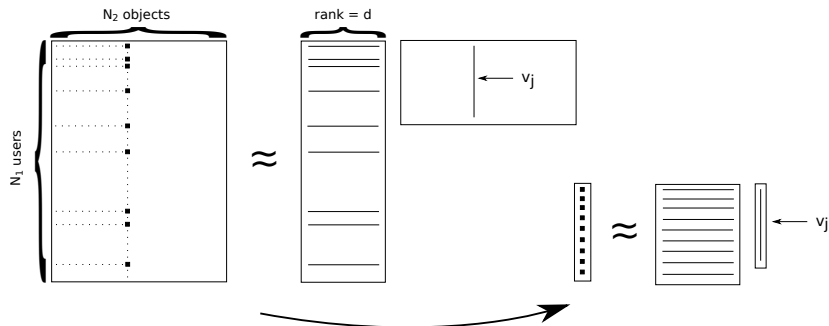
There is a close relationship between this algorithm and ridge regression.

- ▶ Think from the perspective of object location  $v_j$ .
- ▶ Minimize the sum squared error  $\frac{1}{\sigma^2} (M_{ij} - u_i^T v_j)^2$  with penalty  $\lambda \|v_j\|^2$ .
- ▶ This is ridge regression for  $v_j$ , as the update also shows:

$$v_j = \left( \lambda \sigma^2 I + \sum_{i \in \Omega_{v_j}} u_i u_i^T \right)^{-1} \left( \sum_{i \in \Omega_{v_j}} M_{ij} u_i \right)$$

- ▶ So this model is a set of  $N_1 + N_2$  coupled ridge regression problems.

# MATRIX FACTORIZATION AND LEAST SQUARES



We can also connect it to least squares.

- ▶ Remove the Gaussian priors on  $u_i$  and  $v_j$ . The update for, e.g.,  $v_j$  is then

$$v_j = \left( \sum_{i \in \Omega_{v_j}} u_i u_i^T \right)^{-1} \left( \sum_{i \in \Omega_{v_j}} M_{ij} u_i \right)$$

- ▶ This is the least squares solution. It requires that every user has rated at least  $d$  objects and every object is rated by at least  $d$  users.
- ▶ This probably isn't the case, so we see why a prior is *necessary* here.